

E1 Experiment Pipeline

Baseline selection + BERT + ElasticNet regression




Figure 1: E1 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E2 Experiment Pipeline

Baseline selection + BERT + GAN + ElasticNet regression




Figure 2: E2 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E3 Experiment Pipeline

Q-learning selection + BERT + ElasticNet regression




Figure 3: E3 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E4 Experiment Pipeline

Q-learning selection + BERT + GAN + ElasticNet regression




Figure 4: E4 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E5 Experiment Pipeline

Baseline selection + BERT + bidirectional LSTM




Figure 5: E5 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E6 Experiment Pipeline

Baseline selection + BERT + GAN + bidirectional LSTM




Figure 6: E6 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E7 Experiment Pipeline

Q-learning selection + BERT + bidirectional LSTM




Figure 7: E7 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.

E8 Experiment Pipeline

Q-learning selection + BERT + GAN + bidirectional LSTM




Figure 8: E8 experimental training, validation and test pipeline.

Ground-truth Low/High labels use training medians. Decision thresholds use validation data. This is the experimental path; the separate saved-prediction interface embeds the first ten comments per author.