

E1 Experiment Pipeline

Baseline selection, BERT embeddings and ElasticNet regression

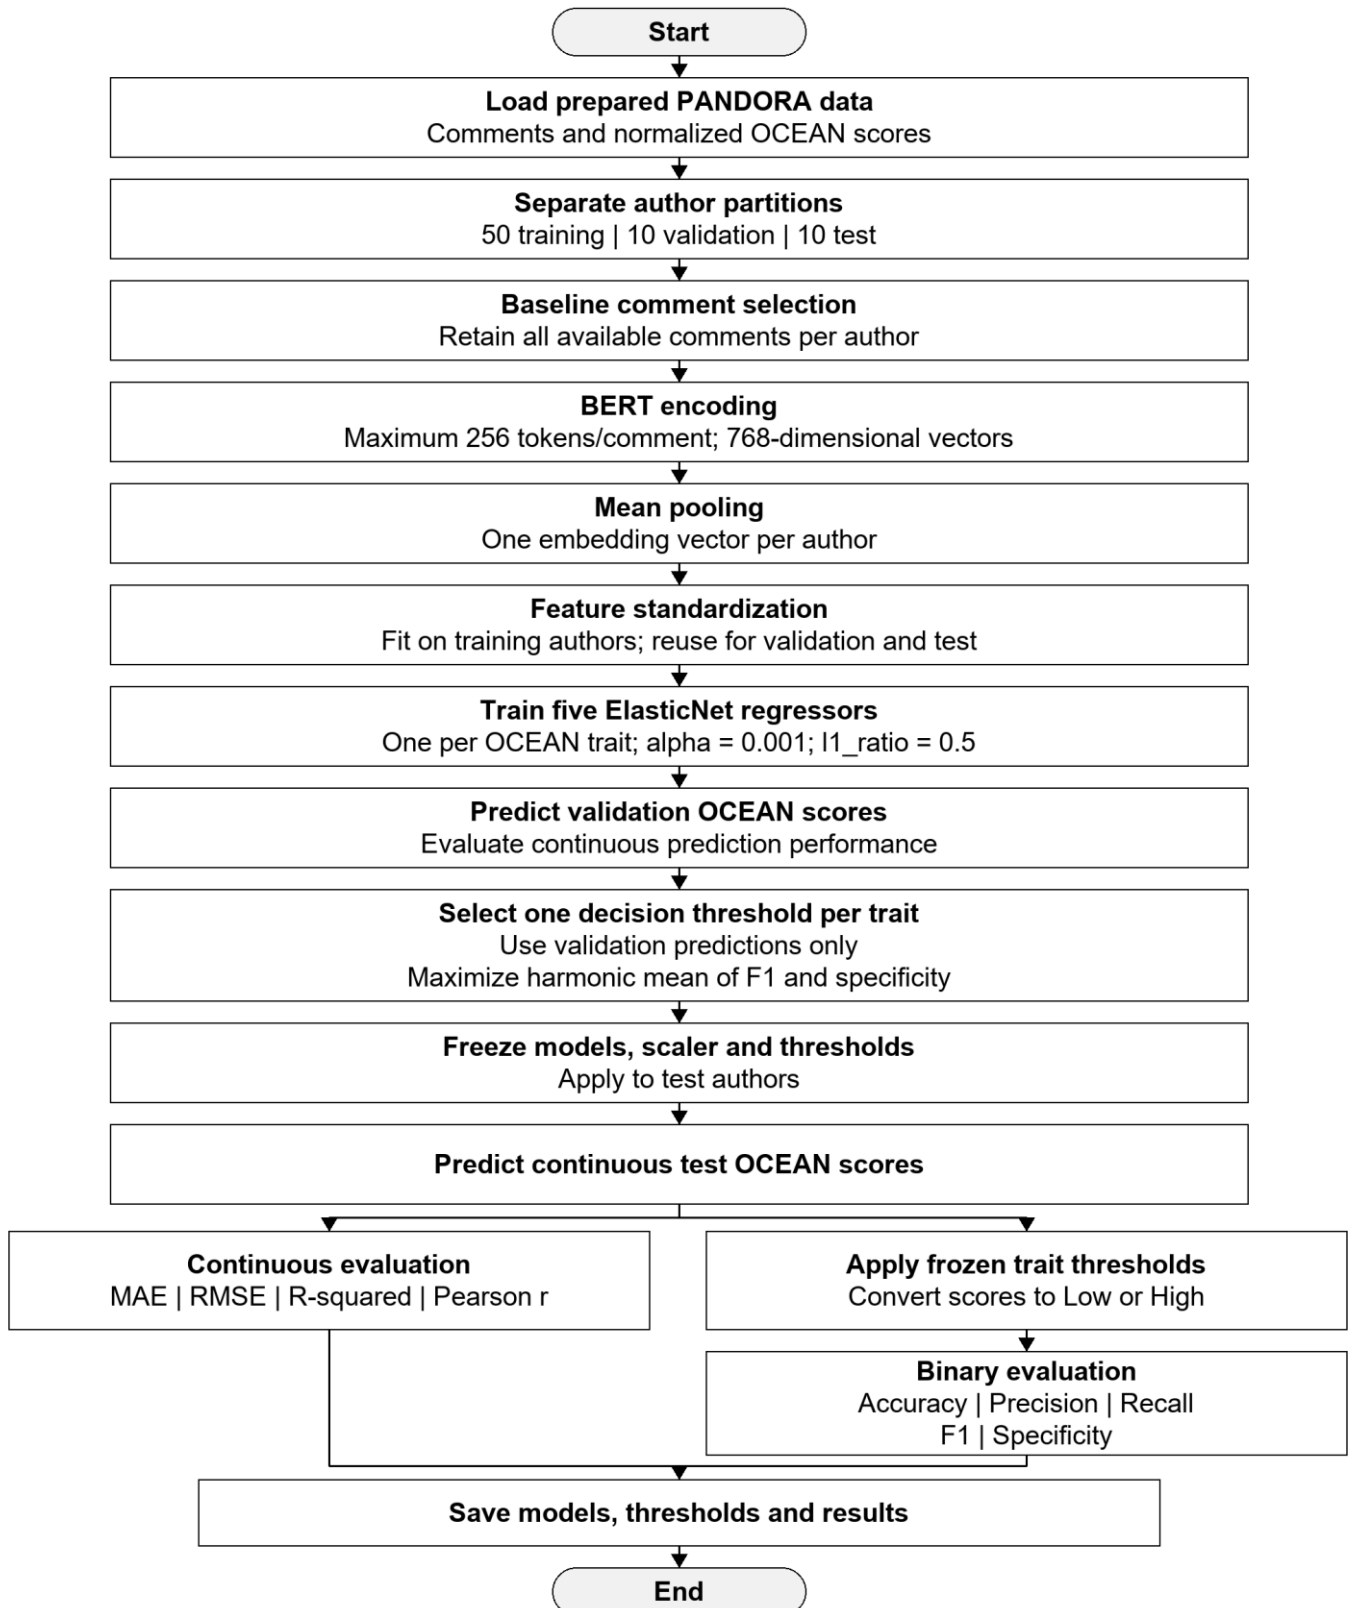


Figure: E1 experimental pipeline. E1 uses no Q-learning or GAN augmentation.

Scope: This figure shows the experimental workflow. The separate saved-prediction interface currently embeds only the first ten comments per author.